

# Recommender Systems

Subash Gandyer  
Data Scientist, HealthChain

# Agenda

- What is a Recommendation Engine?
- Why Recommendations?
- How does a Recommendation Engine work?
- Types of Recommendation
- Examples
  - Simple Recommender
  - Movie Recommender
  - Collaborative Filtering Recommender
  - Matrix Factorization Recommender
- Evaluation Metrics for a Recommendation Engine

# Recommendation Engine

**A recommendation engine filters the data using different algorithms and recommends the most relevant items to users. It first captures the past behaviour of a customer and based on that, recommends products which the users might be likely to buy.**

# Ways of recommending items to users

- Recommend items to a user which are most popular among all the users

Drawback:

Every user will see the **same recommendations**

- Divide the users into multiple segments based on their preferences (user features) and recommend items to them based on the segment they belong to

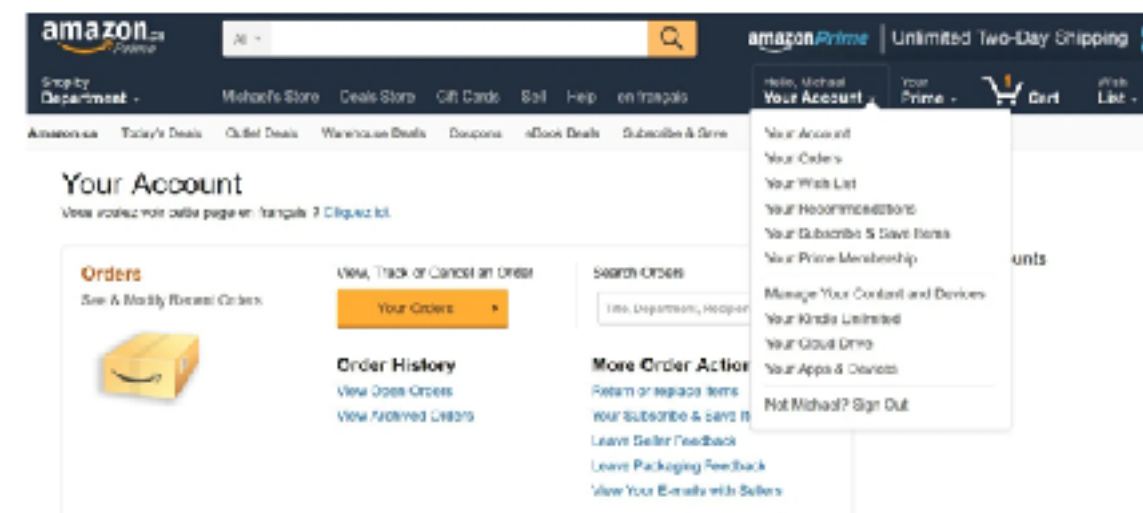
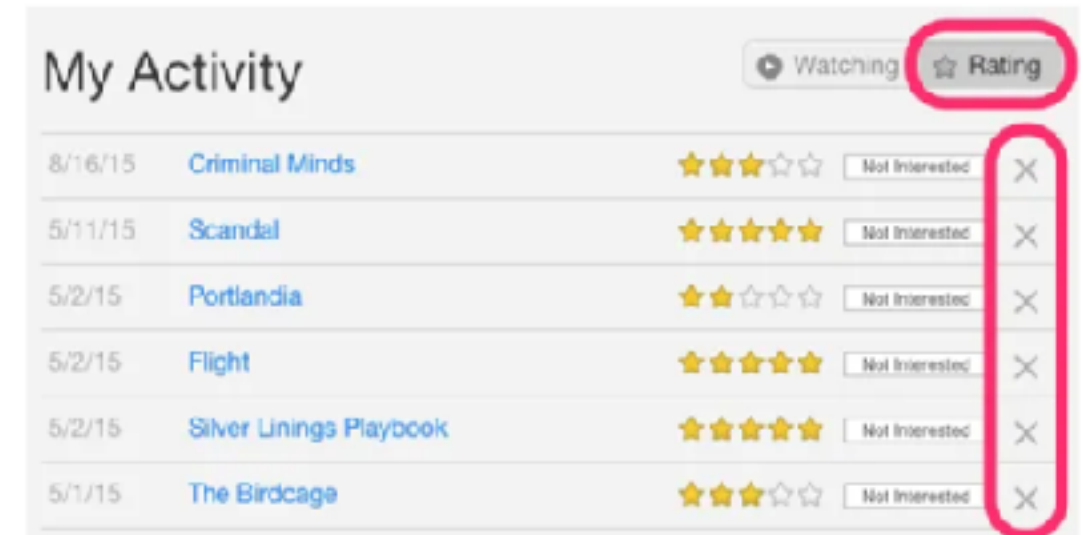
Drawback:

More **computation time** as number of users increases

**Cold Start Problem**

# Data Collection

- Explicit
  - Data provided by users explicitly like Product Ratings, Rankings, etc
- Implicit
  - Data like Clickstream, Search History, Order History, Website Navigation, Web Analytics and so on




# Types of Filtering

- Content-based Filtering
- Collaborative Filtering
- Hybrid

# Content-based filtering

## CONTENT-BASED FILTERING

- Movie recommendation
  - If a person likes “Inception”, then algorithm will recommend movies that fall under the same genre
- How does the algorithm know “**genre**”
  - Profile Vector - Users past behaviour (movies liked/disliked, ratings)
  - Item Vector - movies (genre, cast, director, music, etc)



# Content based filtering (2)

- How recommendations are done?
  - Similarity measured using one of the distances

- Cosine similarity

$$\text{sim}(A, B) = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$

- Euclidean similarity

$$\text{Euclidean Distance} = \sqrt{(x_1 - y_1)^2 + \dots + (x_N - y_N)^2}$$

- Pearsons Correlation

$$\text{sim}(u, v) = \frac{\sum (r_{ui} - \bar{r}_u)(r_{vi} - \bar{r}_v)}{\sqrt{\sum (r_{ui} - \bar{r}_u)^2} \sqrt{\sum (r_{vi} - \bar{r}_v)^2}}$$



# Content based filtering (3)

- Top n
  - Movies are sorted with respect to similarity and top n movies are recommended
- Rating scale
  - Threshold is set and all the movies above threshold are recommended

# Content based filtering (4)

- Drawback:
  - Movies of same genre are recommended
  - Movies of different genres are not recommended

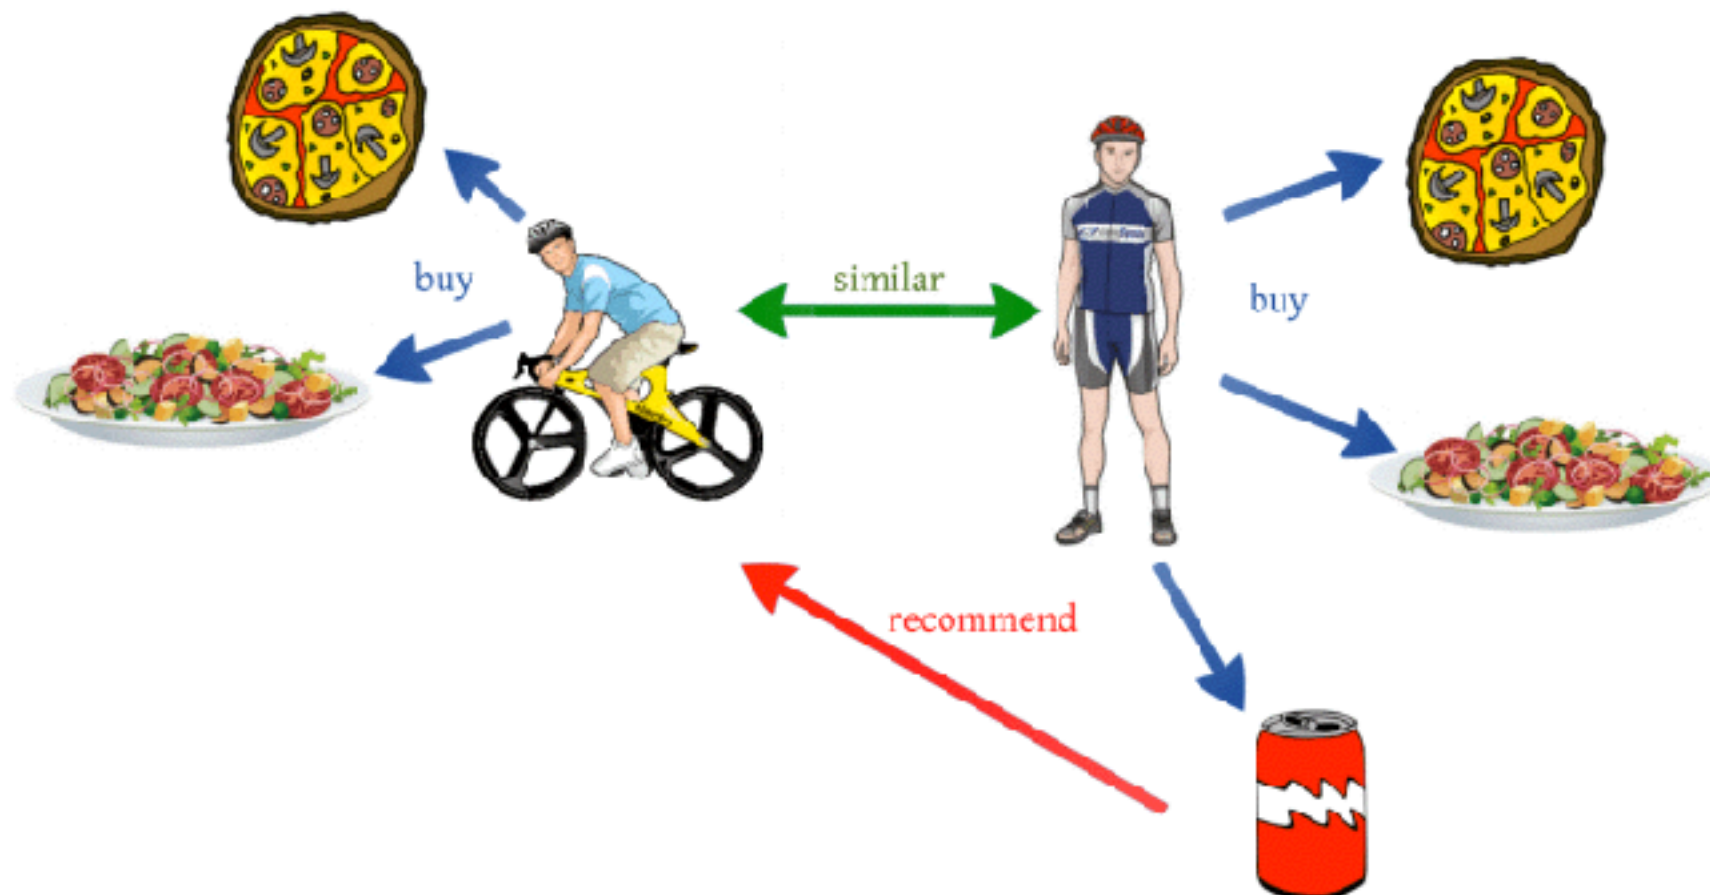
# Collaborative Recommender

# Collaborative filtering

- Algorithm uses **user behaviour** for recommending items
- User-User Collaborative filtering
- Item-Item Collaborative filtering

# User-User Collaborative

- Finds **Similarity score** between users
- For a user, pick the **most similar users**
- **Recommend the items** which these most similar users bought/liked/rated previously



# User-User Collaborative (2)

1. Use Pearson correlation to find similarity between the user  $u$  and  $v$ .
2. First we find the items rated by both the users and based on the ratings, correlation between the users is calculated.
3. The predictions can be calculated using the similarity values. This algorithm, first of all calculates the similarity between each user and then based on each similarity calculates the predictions. **Users having higher correlation will tend to be similar.**
4. Based on these prediction values, recommendations are made

$$P_{u,i} = \frac{\sum_v (r_{v,i} * s_{u,v})}{\sum_v s_{u,v}}$$

# User-User Collaborative (3)

| User/Movie | x1 | x2 | x3 | x4 | x5 | Mean User Rating |
|------------|----|----|----|----|----|------------------|
| A          | 4  | 1  | –  | 4  | –  | 3                |
| B          | –  | 4  | –  | 2  | 3  | 3                |
| C          | –  | 1  | –  | 4  | 4  | 3                |

$$r_{AC} = [(1-3)*(1-3) + (4-3)*(4-3)]/[((1-3)^2 + (4-3)^2)^{1/2} * ((1-3)^2 + (4-3)^2)^{1/2}] = 1$$

$$r_{BC} = [(4-3)*(1-3) + (2-3)*(4-3) + (3-3)*(4-3)]/[((4-3)^2 + (2-3)^2 + (3-3)^2)^{1/2} * ((1-3)^2 + (4-3)^2 + (4-3)^2)^{1/2}] = -0.866$$

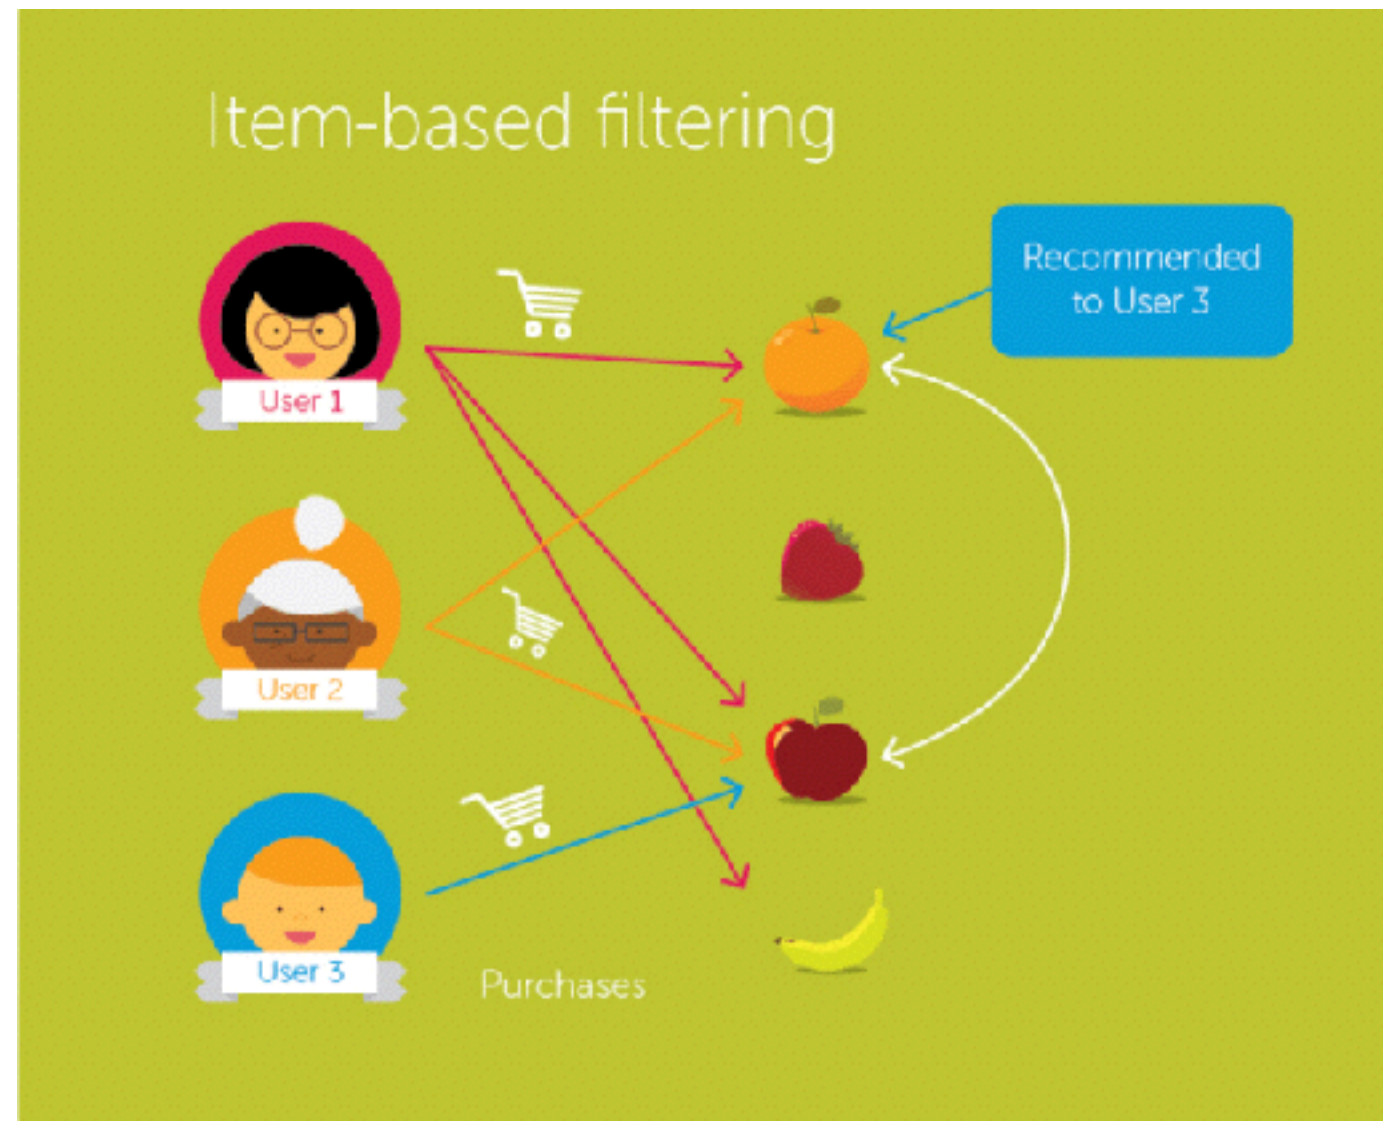
# User-User Collaborative (4)

- Drawbacks:
  - Time consuming
    - Calculates similarity for each user and then calculates prediction for each similarity score
- Solution: (Select few users instead of all)
  - Select a threshold similarity and choose all users above that value
  - Randomly select the users
  - Arrange the neighbours in descending order of their similarity value and choose top-N users
  - Use clustering for choosing neighbours



# Item-Item Collaborative

- Finds **Similarity score** between items
- For an item, pick the **most paired items**
- **Recommend the items** which these most similar users bought/liked/rated previously



# Item-Item Collaborative (2)

| User/Movie       | x1 | x2 | x3 | x4 | x5 |
|------------------|----|----|----|----|----|
| A                | 4  | 1  | 2  | 4  | 4  |
| B                | 2  | 4  | 4  | 2  | 1  |
| C                | –  | 1  | –  | 3  | 4  |
| Mean Item Rating | 3  | 2  | 3  | 3  | 3  |

$$C_{14} = [(4-3)*(4-3) + (2-3)*(2-3)] / [((4-3)^2 + (2-3)^2)^{1/2} * ((4-3)^2 + (2-3)^2)^{1/2}] = 1$$

$$C_{15} = [(4-3)*(4-3) + (2-3)*(1-3)] / [((4-3)^2 + (2-3)^2)^{1/2} * ((4-3)^2 + (1-3)^2)^{1/2}] = 0.94$$

# Item-Item Collaborative (3)

- Movie pairs  $x_1$  and  $x_4$  are highly correlated than  $x_1$  and  $x_5$  pair
- When a user searches for movie  $x_1$ ,  $x_4$  movie is recommended to him and vice versa

# Cold Start problem

- What will happen if a new item or a new user comes into the recommendation ecosystem?
  - No data of that item or user exists
  - No recommendations possible (**Cold Start problem**)
- **Visitor Cold Start** - a new user
  - Solution: **Popularity recommender** - Recommend most popular products to the new user
- **Product Cold Start** - a new item
  - Solution: Content-based filtering - Use content of the new product to recommend

# Movie Recommender

<https://grouplens.org/datasets/movielens/100k/>

# Matrix Factorization Recommender

# Matrix Factorization

Problem:

Not every movie is rated by every user. Find a set of features which can define how a user rates the movies. These are called **latent features**.

We need to find a way to extract the most important latent features from the the existing features.

Solution:

**Matrix factorization** is a technique which uses the lower dimension dense matrix and helps in extracting the important latent features.

# Matrix Factorization (2)

| movie_id | 1   | 2   | 3   | 4   | 5   |
|----------|-----|-----|-----|-----|-----|
| user_id  |     |     |     |     |     |
| 1        | 5.0 | 3.0 | 4.0 | 3.0 | 3.0 |
| 2        | 4.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 3        | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 4        | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 5        | 4.0 | 3.0 | 0.0 | 0.0 | 0.0 |


- User-Item matrix
  - user\_id : every user
  - movie\_id : every item
  - 5.0 - 1.0 (Highly liked - Least liked)
  - 0.0 - Movie not rated by the user



# Matrix Factorization (3)

$$R = P \Sigma Q^T$$

- M is the total number of users
- N is the total number of movies
- K is the total latent features
- R is MxN user-movie rating matrix
- P is MxK user-feature affinity matrix which represents association between users and features
- Q is NxK item-feature relevance matrix which represents association between movies and features
- $\Sigma$  is KxK diagonal feature weight matrix which represents the essential weights of features



# Simple Recommender

# Word2Vec Revisited

we must become what we wish to teach

| Input | Output |
|-------|--------|
| we    | must   |
| we    | become |

we must become what we wish to teach

| Input | Output |
|-------|--------|
| we    | must   |
| we    | become |
| must  | we     |
| must  | become |
| must  | what   |

# Word2Vec Revisited (2)


we must become what we wish to teach

| Input  | Output |
|--------|--------|
| we     | must   |
| we     | become |
| must   | we     |
| must   | become |
| must   | what   |
| become | we     |
| become | must   |
| become | what   |
| become | we     |
| what   | must   |
| what   | become |
| what   | we     |
| what   | wish   |

|       |        |
|-------|--------|
| we    | become |
| we    | what   |
| we    | wish   |
| we    | to     |
| wish  | what   |
| wish  | we     |
| wish  | to     |
| wish  | teach  |
| to    | we     |
| to    | wish   |
| to    | teach  |
| teach | wish   |
| teach | to     |


# Architecture

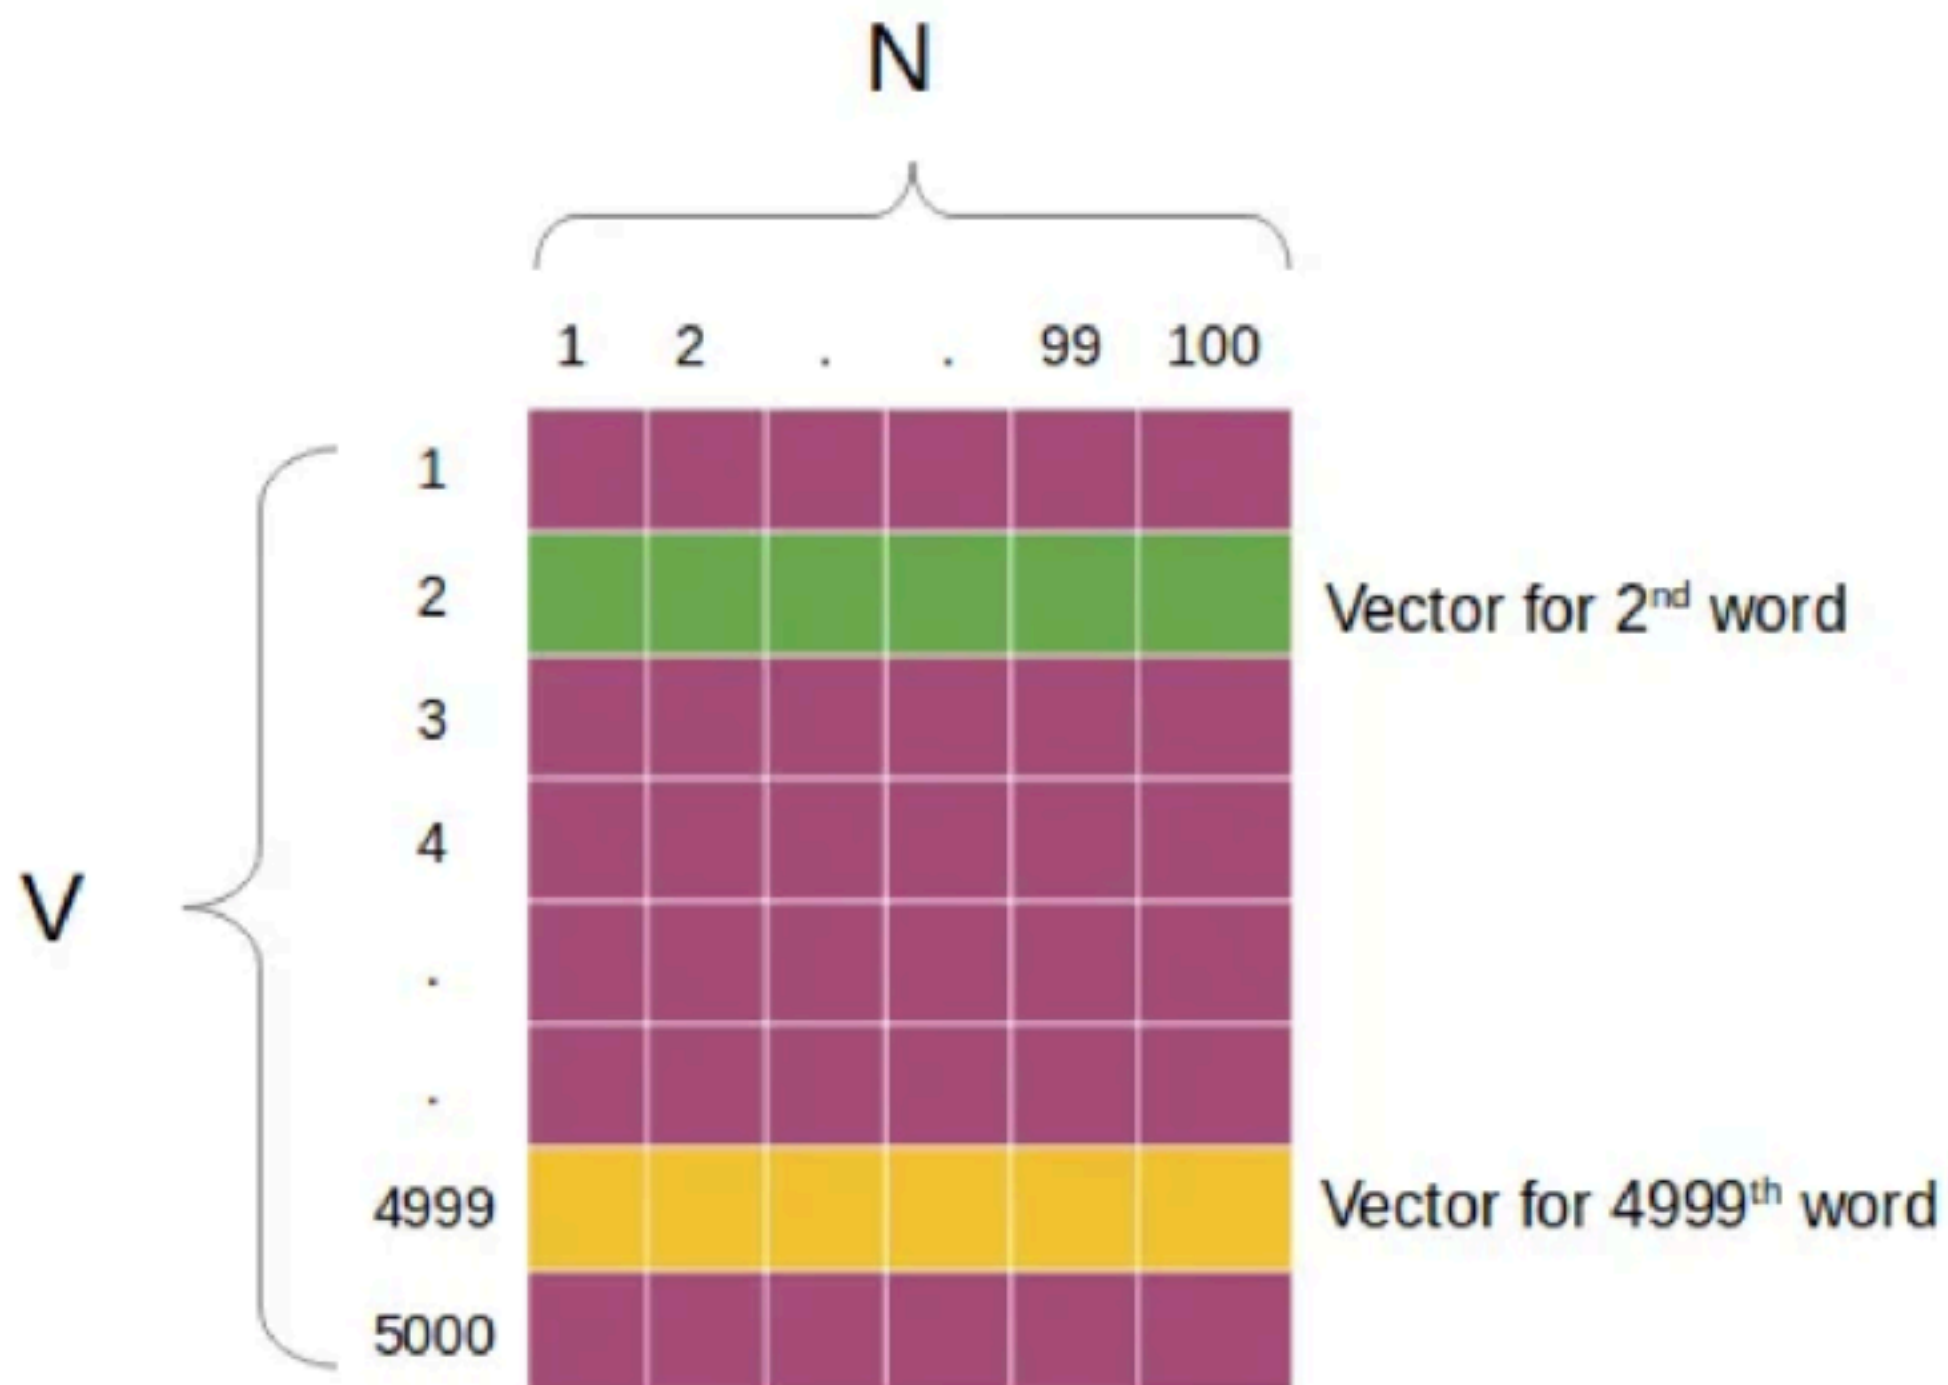
- $V = 5000$  (Size of Vocabulary)
- $N = 100$  (No. Of hidden units / Length of word embeddings)




word2vec model architecture

# Weight Matrix ( $V \times N$ )





# Word Vectors Cosine Similarity





# Problem:

# Product Recommendation



T-shirt



Shorts



Shoes



Cap



Water Bottle

*Purchase history of the consumer*

# Word2Vec Recommendation



# Dataset

**<https://archive.ics.uci.edu/ml/machine-learning-databases/00352/>**

# Lets build a Simple Recommendation System

**Simple\_Recommendation\_System.ipynb**

# Evaluation Metrics

# Evaluation Metrics

- Recall
- Precision
- RMSE
- Mean Reciprocal Rank
- Mean Average Precision
- Normalized Discounted Cumulative Gain

# Recall

- What proportion of items that a user likes were actually recommended

$$\text{Recall} = \frac{tp}{tp + fn}$$

- Here  $tp$  represents the number of items recommended to a user that he/she likes and  $tp+fn$  represents the total items that a user likes
- If a user likes 5 items and the recommendation engine decided to show 3 of them, then the recall will be 0.6
- Larger the recall, better are the recommendations

# Precision

- Out of all the recommended items, how many did the user actually like?

$$\text{Precision} = \frac{tp}{tp + fp}$$

- Here  $tp$  represents the number of items recommended to a user that he/she likes and  $tp+fp$  represents the total items recommended to a user
- If 5 items were recommended to the user out of which he liked 4, then precision will be 0.8
- Larger the precision, better the recommendations
- But consider this case: If we simply recommend all the items, they will definitely cover the items which the user likes. So we have 100% recall! But think about precision for a second. If we recommend say 1000 items and user likes only 10 of them, then precision is 0.1%. This is really low. So, our aim should be to maximize both precision and recall.



# Root Mean Square Error

It measures the error in the predicted ratings

$$RMSE = \sqrt{\frac{\sum_{i=1}^N (Predicted_i - Actual_i)^2}{N}}$$

- Here, Predicted is the rating predicted by the model and Actual is the original rating
- If a user has given a rating of 5 to a movie and we predicted the rating as 4, then RMSE is 1
- Lesser the RMSE value, better the recommendations

Drawback:

No ordering of products recommended is maintained

# Mean Reciprocal Rank

- Evaluates the list of recommendations

$$\text{MRR} = \frac{1}{n} \cdot \sum_{i=1}^n \frac{1}{r(Q_i)}$$

- Suppose we have recommended 3 movies to a user, say A, B, C in the given order, but the user only liked movie C. As the rank of movie C is 3, the reciprocal rank will be 1/3
- Larger the mean reciprocal rank, better the recommendations

# Mean Average Precision

- Precision and Recall don't care about ordering in the recommendations
- Precision at cutoff k is the precision calculated by considering only the subset of your recommendations from rank 1 through k

$$MAP_i = \frac{1}{|R_i|} \sum_{k=1}^{|R_i|} P(R_i[k])$$

- Suppose we have made three recommendations [0, 1, 1]. Here 0 means the recommendation is not correct while 1 means that the recommendation is correct. Then the precision at k will be [0, 1/2, 2/3], and the average precision will be  $(1/3) * (0 + 1/2 + 2/3) = 0.38$
- Larger the mean average precision, more correct will be the recommendations